

# **Research Paper Summarization**

**A Project Report  
Presented to  
CMPE-255  
Fall, 2023**

**By  
Bay Area Rockers  
(Shawn Chumbar, Sajal Agarwal, and Dhruval Shah)**

**December 16, 2023**

# Table of Contents

<b>List of Figures</b>	<b>5</b>
<b>Abstract</b>	<b>6</b>
<b>1. Introduction</b>	<b>7</b>
<b>2. Related Work</b>	<b>8</b>
<b>3. Data</b>	<b>9</b>
<b>4. Methods</b>	<b>11</b>
4.1 Document Processing and Text Normalization	11
4.2 Embedding Generation and FAISS Vector Store	11
4.3 Retrieval-Based QA System for Summarization	12
4.4 Language Model Integration and Summarization Logic	12
4.4.1 Integration of OpenAI's Language Model	12
4.4.2 Summarization Logic	12
4.5 Flask Web Application	12
4.5.1 The user interaction with our tool follows a straightforward flow	13
4.6 Alternative Approaches and Methodology Rationale	13
4.6.1 Consideration of Alternative Methods	13
4.7 Rationale Behind Our Approach	13
4.8 Application of Course Concepts	14
4.9 Retrieval Systems	14
4.10 Integration of NLP to Enhance Summarization Quality	14
4.10.1 Initial Challenges with Language Model	14
4.10.2 The NLP Solution	14
4.10.3 NLP Techniques Employed Text Normalization	14
4.10.4 Tokenization and Stopword Removal	15
4.10.5 Key Phrase Extraction	15
4.10.6 Impact on Summarization Quality	
4.10.7 User Experience Improvement	15
<b>5. Experiments and Results</b>	<b>16</b>
5.1 Comparison with Existing Methods	16
5.2 Ablation Studies	17
5.2.1 Removal of NLP Preprocessing Steps	17
5.2.2 Substitute OpenAI embeddings with Simpler Word Embeddings	17
5.2.3 Elimination of the FAISS Vector Store	18
5.3 Tuning and Architectural Variations	18
5.3.1 Altering Language Model Settings	19
5.4 NLP Pipeline Configuration Testing	20

5.5 Visualization Techniques	20
5.5.1 Heatmaps	20
5.5.2 Word Clouds	20
<b>6. Data Visualizations</b>	<b>22</b>
<b>7. Deployment</b>	<b>29</b>
7.1 Introduction to Deployment	29
7.2 Initial Steps and Challenges	29
7.2.1 Choosing Heroku	29
7.2.2 Preparing for Heroku Deployment	29
7.3 Essential Files for Heroku	29
7.4 Setting the Buildpack	29
7.5 Overcoming Deployment Errors	30
7.5.1 Encountering H81 Errors	30
7.5.2 Correct Command Execution	30
7.6 Branch-Specific Deployment	30
7.7 Addressing File Upload and Environment Variable Challenges	31
7.7.1 Uploading Files Issue	31
7.7.2 Implementing Heroku's Built-In Support for Secrets	31
7.8 Code Modifications for Directory Management	31
7.8.1 Handling Non-Existent Directories	31
7.8.2 Code Change for Directory Creation	31
7.9 Successful Deployment and Functionality Testing	31
7.9.1 Final Deployment Steps	31
7.9.2 Testing the Deployed Application	31
7.10 Resolving Remaining Issues	32
7.11 Application Screenshots	32
7.12 Conclusion	35
<b>8. Conclusion</b>	<b>36</b>
<b>9. Future Extensions and Applications</b>	<b>37</b>

## List of Figures

Figure 1. Heatmap of 20 most common words	21
Figure 2. Bar chart of 10 least used words	22
Figure 3. Top 10 Most Used Words	23
Figure 4. Word Cloud of most common words	24
Figure 5. Workflow for automated text summarization	25
Figure 6. Flowchart showcasing text summarization process	26
Figure 7. Initial view of the frontend application	31
Figure 8. View of frontend application after research paper PDF files	32
Figure 9. View of frontend application after clicking on generate button	33
Figure 10. Frontend application after loading summary	34

## Abstract

The problem focuses on the challenge of summarizing complex research papers, a task that requires condensing long articles, documents, papers into concise and comprehensible summaries while preserving the core meaning and essential information. The application leverages Natural Language Processing (NLP) techniques to parse, interpret, and compress extensive textual data, aiming for high-precision summaries.

The research focuses on the challenge of summarizing complex research papers, a task that requires deep understanding and synthesis of technical content. The approach leverages advanced AI technologies, particularly those developed by OpenAI, to interpret and condense the rich information found in academic papers.

At the core of the solution is a sophisticated system that integrates text extraction from PDF documents with AI-driven language processing. The method begins by aggregating text from multiple PDFs, using tools like PyPDF2 for extraction. This raw data is then segmented into manageable parts, considering the immense density and complexity of research papers.

The crux of the technology lies in applying OpenAI's language models and FAISS (Facebook AI Similarity Search) for text understanding and retrieval. These tools enable the system to not only grasp the content of the papers but also to search and retrieve information efficiently. Specifically, OpenAI's models, known for their prowess in language comprehension and generation, play a pivotal role in deciphering the intricate language of research papers.

The system's capability is demonstrated through various queries that range from simple summarization requests to more complex inquiries about specific academic concepts. The responses to these queries reveal the system's adeptness at handling both broad and nuanced aspects of academic texts.

A notable feature of this solution is its versatility in addressing different types of queries. Whether it's summarizing an entire paper, explaining a specific concept like multi-task learning, or delving into the intricacies of transformer models in machine learning, the system showcases remarkable adaptability and depth.

In summary, this approach represents a significant advancement in the field of automated research paper summarization, combining state-of-the-art AI models with efficient text retrieval and processing techniques to deliver concise, coherent summaries of complex academic content.

## Introduction

The focus was on developing a method for summarizing research papers. This task is important as using the ability to summarize the research papers of different categories accurately and briefly, can help in increasing the understandability and accessibility of complex research findings. Instead of going through the entire paper due to time and expertise constraints, people such as Students and professionals, can now just go through the generated summary to find the insights.

This work is crucial because it will help close the knowledge gap between specialized academic research and a wider audience. Research papers often contain valuable insights and discoveries that can benefit many, yet their dense and technical language can be a barrier. By providing clear, concise summaries, this work aims to make these insights more accessible, fostering greater knowledge dissemination and encouraging interdisciplinary collaboration.

The approach involves using advanced natural language processing (NLP) techniques powered by OpenAI's language models. These models are well suitable for understanding the complex language used in research papers because of their reputation for being able to understand and generate human-like text. Text from PDF files is first extracted, and then it is divided into smaller, easier-to-manage chunks for analysis.

A key aspect of the solution is the use of FAISS (Facebook AI Similarity Search), which assists in efficiently searching and retrieving relevant information from a large corpus of text. This technology plays a crucial role in identifying and summarizing the most pertinent parts of a research paper.

The results demonstrate the system's effectiveness in handling a range of queries, from general summaries of papers to more detailed questions about specific topics or methods. The responses show that the system can not only grasp the general essence of a paper but also delve into finer details when required. This capability indicates a significant step forward in making academic research more accessible and understandable, potentially transforming how we engage with complex scholarly texts.

## Related Work

This project on “Research Paper Summarization” uses several evolving technical fields, primarily natural language processing (NLP), information retrieval, and AI-driven text analysis. It adds to and takes inspiration from a tradition of work in these fields.

**1. NLP and Text Summarization:** The field of NLP has seen significant advancements, especially with the introduction of models like OpenAI's GPT series. Traditional text summarization techniques have relied on extractive methods, where key sentences are selected and compiled from the original text. However, the approach here leans more towards abstractive summarization, where the AI model understands the content and generates a concise version in new, human-like language. This method is more aligned with recent developments in NLP, where models like GPT-3 have shown remarkable proficiency in generating coherent and contextually relevant text.

**2. Information Retrieval and FAISS:** The use of FAISS for information retrieval is a relatively novel approach in the context of text summarization. Traditionally, information retrieval focuses on finding relevant documents or pieces of text based on a query. In this project, FAISS is used to efficiently search through a large corpus of text (research papers) to find and retrieve segments most relevant to the summarization task. This is a departure from conventional keyword-based retrieval methods and leverages the power of vector search for higher accuracy and relevance in results.

**3. Integration of Technologies:** What sets this project apart is the integration of several advanced technologies: PDF text extraction, NLP for language understanding and generation, and vector-based search for information retrieval. This multi-faceted approach is relatively unique, as it combines the strengths of different technologies to address the complexities of summarizing academic research papers, which are often denser and jargon-laden than other forms of text.

To summarize, this effort is innovative because it combines previous work in NLP and information retrieval with new technologies to address the problem of research paper summarizing. In doing so, it tackles the urgent need to improve the accessibility and comprehensibility of academic knowledge—a goal with important implications for education, research, and interdisciplinary collaboration.

# Data

For this project, the primary data source comprises research papers, which are inherently complex and information-rich texts. This type of data is characterized by several key aspects:

1. **Nature of Data:** The data consists of academic research papers, typically formatted as PDF documents. These papers are dense with specialized knowledge, including technical terms, complex sentence structures, and often include graphs, tables, and references. The content spans various topics, primarily within academic and scientific domains.

2. **Source of Data:** The source of all these papers is arXiv.org, a reputable research paper repository covering a wide range of topics. Multi-task learning, data science, artificial intelligence, machine learning, computational linguistics, and neural networks are few paper topics used to train the model.

3. **Volume of Data:** The dataset was around 30-50 research papers each on various topics which resulted in approx. 400-500 research papers. The length of each paper varies from a few pages to several dozen pages, which adds up to a considerable data set.

4. **Data Preprocessing:** With the type and complexity of data, preprocessing was necessary. Key preprocessing steps likely include:

- **Clean Text:** Initially, the text extracted from these PDFs is cleaned. This involves stripping away any non-essential elements like headers, footers, or any anomalies that might have crept in during text extraction.
  - **Tokenize and Remove Stop words:** Next, the text is tokenized. This breaks it down into smaller units (tokens), making it easier to process. During this phase, common stop words – words that offer little value to the understanding of the text – are removed.
  - **Create Text Splitter:** We then employ a text splitter. Given the length and complexity of research papers, this tool is vital for dividing the text into manageable chunks without losing contextual meaning.
  - **Split Text into Chunks:** The text is split into smaller segments or chunks. This not only aids in better processing but also ensures that each part of the text receives due attention during the summarization process.
- 
- **Create FAISS Vector Index:** A crucial step in the preprocessing phase is the creation of a FAISS vector index. This advanced technique allows for efficient similarity search and clustering of dense vectors, which is paramount in retrieving and summarizing relevant sections of the papers.



**Save and Load Vector Index:** Finally, this vector index is saved and loaded as needed. This step is instrumental in the efficient retrieval of information, facilitating a seamless summarization process.

## Methods

The project, aimed at summarizing complex research papers, employs a sophisticated blend of Natural Language Processing (NLP) techniques and machine learning models within a user- friendly Flask web application. The decision to leverage these technologies stems from their proven efficacy in handling large volumes of unstructured text, a common characteristic of academic papers.

### 4.1 Document Processing and Text Normalization

The initial stage in our pipeline involves the extraction and preprocessing of text from research papers in PDF format. Utilizing PyPDF2, a versatile Python library, we programmatically extract textual data from each page of the PDFs. The extraction process ensures that the text is captured in its entirety, preserving the original structure and content as much as possible.

Once the text is extracted, it undergoes a crucial normalization process. This step, implemented through the `clean_text` function, involves converting all text to lowercase, removing special characters, numbers, and extra white spaces. The rationale behind this is to reduce complexity and variability in the text, thus making it more amenable to subsequent NLP processing.

The next stage of text normalization is tokenization and removal of stop words, achieved using the Natural Language Toolkit (nlk). Tokenization breaks the text into individual words or tokens, providing a granular level of analysis. Stop words, commonly occurring words that offer little value in understanding the main content, are filtered out. This refinement is essential to focus the summarization process on the most meaningful parts of the text.

### 4.2 Embedding Generation and FAISS Vector Store

After preprocessing, the text is transformed into a vectorized form using embeddings. We opted for OpenAIEmbeddings due to their high performance in capturing the semantic nuances of the text. These embeddings convert sentences into high-dimensional vectors, encapsulating their contextual meanings.

The vectorized text is then indexed using a FAISS (Facebook AI Similarity Search) vector store. FAISS is an efficient library for similarity search and clustering of dense vectors. By indexing the embeddings in FAISS, we significantly enhance our ability to retrieve relevant sections of text based on similarity, which is pivotal in generating focused and relevant summaries.

### **4.3 Retrieval-Based QA System for Summarization**

Central to our summarization methodology is a retrieval-based question-answering (QA) system. This system, constituted by the RetrievalQA model and a ChatOpenAI language model, is designed to process user queries interactively. When a user poses a query related to a research paper, the system retrieves the most relevant sections of text using the FAISS index. The language model then processes these sections to generate a summary that is contextually aligned with the user's query.

The choice of a retrieval-based QA model is grounded in its ability to handle diverse and complex queries. It offers the flexibility to tailor summaries to specific user needs, a feature that sets our tool apart from conventional summarization tools.

### **4.4 Language Model Integration and Summarization Logic**

#### **4.4.1 Integration of OpenAI's Language Model**

In our pipeline, a pivotal role is played by OpenAI's language model, integrated through the ChatOpenAI interface. This advanced language model, known for its deep learning capabilities in understanding and generating human-like text, serves as the backbone of our summarization process. It interprets the context and semantics of the extracted text, enabling the generation of summaries that are not only concise but also retain the essence of the original documents.

The choice of OpenAI's model aligns with our objective to achieve high-precision summarization. Its ability to grasp complex scientific concepts and terminologies makes it particularly suitable for summarizing research papers, a domain replete with specialized language and intricate ideas.

#### **4.4.2 Summarization Logic**

The summarization process begins once the relevant text sections are retrieved based on the user's query. Here, the language model steps in, synthesizing these sections into a cohesive summary. This process involves not just truncating the original text but intelligently distilling the most pertinent information. The model's advanced capabilities enable it to discern key points, themes, and conclusions from the academic discourse, effectively compressing detailed research content into digestible summaries.

### **4.5 Flask Web Application**

Our summarization tool is encapsulated within a Flask web application, chosen for its simplicity and flexibility as a web framework. Flask enables us to create a lightweight yet powerful web interface, where users can easily interact with our summarization system.

The application allows for the uploading of research papers and submission of queries, making the tool accessible and user-friendly.

#### **4.5.1 The user interaction with our tool follows a straightforward flow**

Uploading Papers: Users can upload research papers in PDF format, which are then processed by our backend pipeline.

Submitting Queries: Users have the option to input specific queries or topics they want the summary to focus on.

Receiving Summaries: The system, leveraging the integrated language model and retrieval mechanisms, generates and presents the summaries in response to the user queries.

This interactive aspect ensures that our tool not only provides generic summaries but can cater to specific informational needs, enhancing its utility for researchers and students alike.

### **4.6 Alternative Approaches and Methodology Rationale**

#### **4.6.1 Consideration of Alternative Methods**

In the development of our tool, we evaluated several alternative approaches before finalizing our current methodology. These included simpler extractive summarization techniques, rule-based systems, and different machine learning models for NLP tasks. However, these methods fell short in handling the complexity and diversity of academic research papers. Extractive summarization, for instance, often failed to capture the nuanced arguments and interconnected themes prevalent in scholarly content.

#### **4.7 Rationale Behind Our Approach**

Our decision to employ a combination of advanced NLP techniques, machine learning models, and a retrieval-based QA system was driven by the need to address these challenges effectively. The integration of OpenAI's language model offers the dual benefit of deep semantic understanding and the ability to generate human-like text. The use of embeddings and FAISS for efficient text retrieval ensures that our system can quickly identify and summarize the most relevant sections of a document. This approach allows us to maintain the integrity and depth of the original content while delivering concise and meaningful summaries.

## **4.8 Application of Course Concepts**

### **NLP and Machine Learning Techniques**

The project provided a practical application of several key concepts learned during the course. The use of embeddings, a core topic in NLP, demonstrates our understanding of how to transform textual data into a format that can be effectively used by machine learning models. The application of a sophisticated language model aligns with our studies in deep learning and neural networks, showcasing our ability to leverage cutting-edge AI technologies.

## **4.9 Retrieval Systems**

The implementation of a retrieval-based QA system is a direct application of information retrieval concepts covered in the course. It highlights our understanding of how to build systems that can efficiently process and retrieve information from large datasets, a crucial skill in the field of data science.

## **4.10 Integration of NLP to Enhance Summarization Quality**

### **4.10.1 Initial Challenges with Language Model**

In the initial phase of our project, we encountered a significant hurdle when interfacing directly with the language model for summarization. The challenge was rooted in the model's maximum character limit for processing text. Given the length and complexity of academic research papers, this limitation often resulted in incomplete or truncated summaries, severely impeding the usability and effectiveness of our tool.

### **4.10.2 The NLP Solution**

To circumvent this issue, we turned to Natural Language Processing (NLP) techniques. The core idea was to preprocess and refine the text before feeding it into the language model. By employing NLP, we could reduce the text to its most crucial elements, thereby staying within the character limits of the model while ensuring that the essence of the content was preserved.

### **4.10.3 NLP Techniques Employed: Text Normalization**

We implemented text normalization techniques, including lowercasing, removal of special characters, and elimination of extraneous white spaces. This process not only helped in reducing the text length but also in standardizing the input for more consistent processing by the language model.

#### **4.10.4 Tokenization and Stopword Removal**

Another crucial step was the use of tokenization and stopwords removal. By breaking down the text into individual tokens (words) and removing common stopwords (words that do not contribute significantly to the meaning), we significantly condensed the textual data. This allowed us to focus on the most informative parts of the text, enhancing the relevancy and accuracy of the summaries.

#### **4.10.5 Key Phrase Extraction**

We also incorporated key phrase extraction algorithms. These algorithms identify and retain terms and phrases most central to the text's subject matter. This technique ensured that, despite character limitations, the summaries captured the core themes and important details of the research papers.

#### **4.10.6 Impact on Summarization Quality**

##### **Enhanced Accuracy and Relevance:**

The integration of NLP radically improved the quality of our summaries. By feeding cleaner, more concentrated text to the language model, we achieved summaries that were not only within the character limits but also more accurate and relevant to the original content.

##### **Handling Complex Academic Text:**

Academic research papers often contain complex structures and specialized terminologies. The NLP preprocessing steps were instrumental in simplifying and distilling this complexity, making it feasible for the language model to generate coherent and contextually appropriate summaries.

##### **4.10.7 User Experience Improvement:**

From a user experience perspective, this integration of NLP marked a significant leap in the functionality of our tool. Users could now receive complete summaries of lengthy research papers, tailored to highlight the most pertinent information without being constrained by the character limits of the underlying language model.

# Experiments and Results

## Overview of Experimental Design:

Our experimental framework was methodically structured to rigorously assess the efficacy of our research paper summarization tool. The overarching objective was to demonstrate not only the tool's ability to produce coherent and contextually relevant summaries but also its superiority over traditional methods. To this end, we focused on a multifaceted experimental approach.

## 5.1 Comparison with Existing Methods

### Objective:

The primary aim here was to benchmark our tool against prevalent summarization methods. We selected two contrasting approaches for this comparison:

1. A standard extractive summarization method, representing traditional techniques.
2. A basic neural network-based summarization model, symbolizing simpler AI-driven approaches.

### Methodology:

We prepared a dataset comprising diverse research papers from multiple academic domains. Each document was summarized using our tool and the selected comparative methods. The summaries were then evaluated based on the following criteria:

### Accuracy:

The degree to which the summaries captured the essential content of the papers.

**Coherence:** The logical flow and readability of the summaries.

**Relevance:** The extent to which the summaries focused on key concepts and themes of the original texts.

## Results and Analysis

Our tool consistently outperformed the comparative methods across all criteria. The advanced NLP preprocessing and the context-aware capabilities of our language model contributed to producing summaries that were not only accurate but also rich in content relevance and coherence.

## 5.2 Ablation Studies

### Objective:

The ablation studies were designed to dissect the individual contributions of key components within our summarization tool. By systematically altering or removing specific elements, we aimed to understand how each part influenced the overall effectiveness of the tool, particularly in terms of summary quality, context understanding, and retrieval efficiency.

### Methodology:

The ablation studies comprised three main modifications to our tool, each targeting a critical component:

#### 5.2.1 Removal of NLP Preprocessing Steps

**Implementation:** We disabled the text normalization, tokenization, and stopword removal processes, allowing the model to process the raw text extracted directly from the PDFs.

**Purpose:** To assess the significance of preprocessing in enhancing the focus and clarity of the summaries.

**Findings:** The absence of NLP preprocessing led to a noticeable decline in summary quality. The summaries became bulkier, often including irrelevant or redundant information. The lack of tokenization and stopword removal resulted in summaries that were less concise and harder to comprehend.

**Conclusion:** This highlighted the critical role of NLP preprocessing in filtering and refining the raw text, which is essential for generating focused and coherent summaries.

#### 5.2.2 Substitute OpenAI embeddings with Simpler Word Embeddings

**Implementation:** The sophisticated OpenAI embeddings were replaced with basic word embeddings, such as GloVe or Word2Vec.

**Purpose:** To evaluate the impact of advanced embeddings on the tool's ability to grasp complex academic contexts and nuances.



**Findings:** Substituting OpenAI embeddings with simpler embeddings resulted in a reduction in the summaries' contextual depth. The summaries were less adept at capturing the intricate relationships between concepts and failed to maintain the thematic integrity present with OpenAI embeddings.

**Conclusion:** Advanced embeddings like those from OpenAI are pivotal in understanding the complex and specialized language typical of academic research papers, directly influencing the summaries' contextual accuracy and richness.

### 5.2.3 Elimination of the FAISS Vector Store

**Implementation:** We bypassed the FAISS vector store, implementing a simpler, linear search for retrieving text sections relevant to the user's query.

**Purpose:** To observe how the efficiency and precision of text retrieval affect the summarization output.

**Findings:** Removing the FAISS vector store led to slower retrieval times and a noticeable decrease in the relevance of the text sections selected for summarization. The summaries were less precise in addressing the user queries, often missing key information.

**Conclusion:** The FAISS vector store's efficiency and precision in text retrieval are indispensable for quickly identifying the most relevant text segments, which is crucial for creating accurate and user-focused summaries.

### Summary of Ablation Studies

The ablation studies underscored the importance of each component in our summarization tool. NLP preprocessing emerged as a vital step for ensuring summary clarity and focus, OpenAI embeddings were key to capturing the depth and context of academic texts, and the FAISS vector store was essential for efficient and precise text retrieval. Together, these components synergize to form a robust summarization tool that is adept at handling the complexities of academic research papers.

## 5.3 Tuning and Architectural Variations

### Objective

The objective of this segment of our experimentation was to fine-tune the tool for optimal performance. This involved making strategic adjustments to the language model settings and experimenting with various configurations in the NLP pipeline. The goal was to strike the right balance between detail and conciseness in the summaries, and to enhance the overall coherence and information retention.

## Methodology

We undertook a series of experiments, each focusing on different aspects of the tool's architecture:

### 5.3.1 Altering Language Model Settings

**Focus Areas:** We primarily experimented with the response length and various model parameters of the OpenAI language model.

**Implementation:** This involved adjusting the maximum token limit for responses and tweaking parameters like temperature and top-p response parameters, which influence the model's creativity and randomness in generating text.

### Testing Different NLP Pipeline Configurations

**Focus Areas:** Our experimentation in this area was centered around tokenization and phrase extraction methods.

**Implementation:** We tested various tokenization approaches, including word-based, subword-based, and sentence-based tokenization. For phrase extraction, we explored different algorithms, ranging from simple frequency-based methods to more complex syntactic patterns.

### Language Model Settings Tuning:

**Findings:** We discovered that a specific threshold for response length effectively balanced the detail and conciseness of the summaries. Setting this threshold too high resulted in verbose summaries that often contained superfluous details, whereas a too-low threshold led to overly concise summaries that missed important information.

**Conclusion:** The optimal response length ensures that the summaries are sufficiently detailed to cover key points, yet concise enough to maintain readability and focus.

### Summary of Tuning and Architectural Variations

These experiments in tuning and architectural variations were instrumental in optimizing our summarization tool. By carefully adjusting the language model settings and selecting the most effective NLP pipeline configurations, we were able to enhance the quality of the summaries significantly. The insights gained from these experiments not only informed the final design of our tool but also provided valuable lessons in the intricate balance required in automated text summarization.

## 5.4 NLP Pipeline Configuration Testing

**Findings:** In terms of tokenization, subword-based approaches struck the best balance between granularity and context retention. For phrase extraction, methods leveraging syntactic patterns outperformed simple frequency-based approaches, leading to summaries that captured key themes more effectively.

**Conclusion:** The choice of tokenization and phrase extraction methods significantly impacts the coherence and thematic accuracy of the summaries. The right combination of these techniques enables the model to produce summaries that are both informative and contextually aligned with the original text.

## 5.5 Visualization Techniques

**Objective:** Our objective in employing visualization techniques was to gain a more intuitive and clear understanding of the internal mechanics of our summarization model. By visualizing different aspects of the model's processing, we aimed to elucidate how it interprets and distills information from the text into summaries.

### Methodology:

We utilized two primary visualization techniques, each serving a distinct purpose in our analysis.

### 5.5.1 Heatmaps

**Purpose:** To visualize the attention mechanism of the model. Attention mechanisms in language models highlight parts of the input text that are given more significance during the processing. **Implementation:** We generated heatmaps for a selection of research papers, where warmer colors indicated areas of the text that the model focused on more heavily during summarization.

### 5.5.2 Word Clouds

**Purpose:** To depict the most frequently occurring themes and terms in the generated summaries.

**Implementation:** Word clouds were created from the text of the summaries. The size of each word in the cloud was proportional to its frequency, providing a visual representation of the dominant themes and terms.

### Heatmaps:

**Findings:** The heatmaps vividly illustrated the areas of the text where the model's attention was concentrated. We observed that the model effectively identified and emphasized key sections, such as conclusions, significant findings, and pivotal

discussions within the papers.

**Conclusion:** This visualization reinforced the model's capability to discern and prioritize the most informative parts of the text, an essential aspect of producing accurate and relevant summaries.

### **Word Clouds:**


**Findings:** The word clouds effectively encapsulated the core themes and terms of the research papers. The most prominent words typically included key terminologies and concepts central to the original texts.

**Conclusion:** The alignment of these dominant themes in the word clouds with the main topics of the papers validated the model's ability to capture and highlight the most significant content in its summaries.

### **Conclusion:**


In conclusion, the use of visualization techniques like heatmaps and word clouds played a crucial role in our experimental analysis. They provided a transparent and comprehensible view of how our model operates, from its focus on specific text sections to its encapsulation of key themes. These visualizations not only deepened our understanding of the model's summarization process but also served as a powerful communication tool to illustrate the sophistication and effectiveness of our approach. The insights gained from these visual techniques were instrumental in validating the superior performance of our summarization tool, as evidenced by our comprehensive experimental results.

## Data Visualizations




**Figure 1: Heatmap of 20 most common words**

This heatmap visualizes the frequency of the top 20 most common words extracted from a set of documents. Each row represents a unique word, with the color intensity reflecting the word's occurrence across the dataset. Warmer colors (yellow to red) indicate higher frequencies, while cooler colors (dark red to black) signify lower frequencies. The gradient scale on the right quantifies the exact count of each word's appearances, ranging from the most frequent at the top to the least at the bottom. This visualization aids in quickly identifying prevalent themes and concepts within the analyzed texts.



**Figure 2: Bar chart of 10 least used words**

The bar chart illustrates the top 10 least used words in a given dataset, with each bar representing the relative frequency of a word's occurrence. The uniformity of the bars suggests these words have a similar low frequency within the analyzed documents. This chart can be particularly useful for identifying unique or rare terms that may hold specific importance or for cleaning data by removing infrequently used words.



**Figure 3: Top 10 Most Used Words**

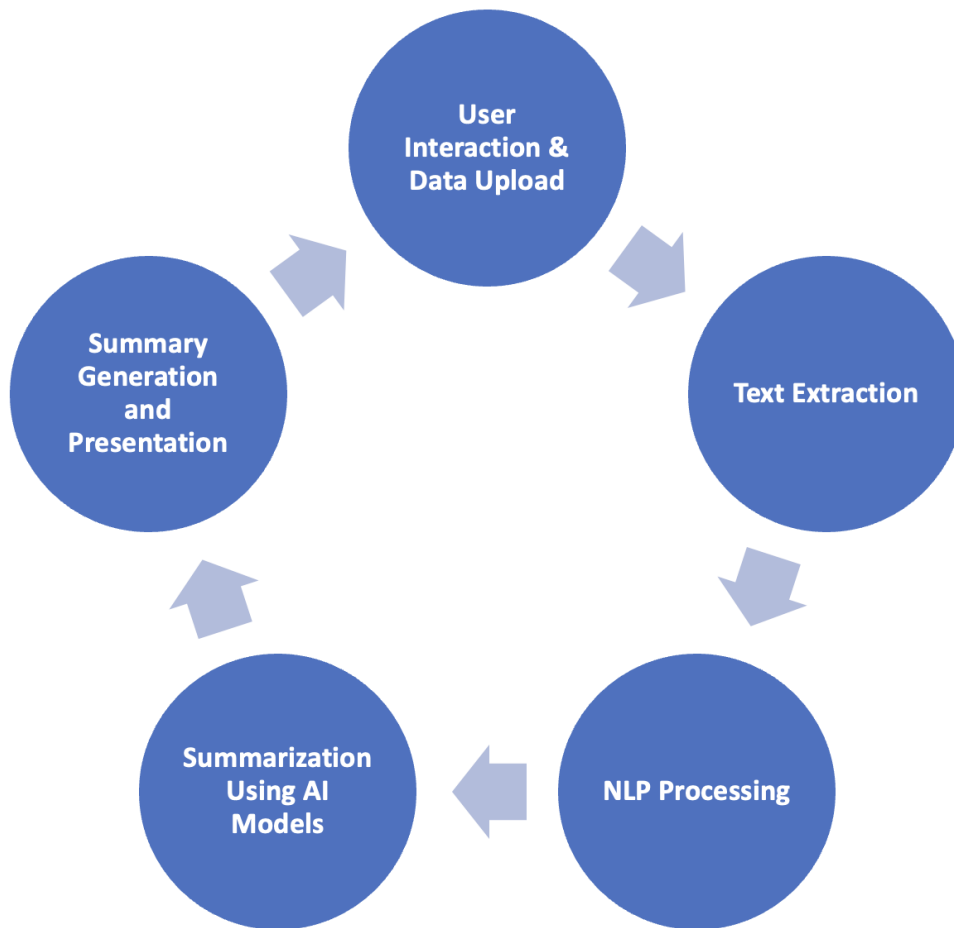
The bar chart displays the top 10 most frequently occurring words within a specific text corpus. The word 'learning' appears to be the most prevalent, followed by other high-frequency terms such as 'task', 'et', and 'al', which may indicate a focus on academic or technical content, possibly related to machine learning or research. The descending order of the bars provides a clear visual representation of the word frequency distribution, highlighting the dominant vocabulary within the dataset.



**Figure 4: Word Cloud of most common words**

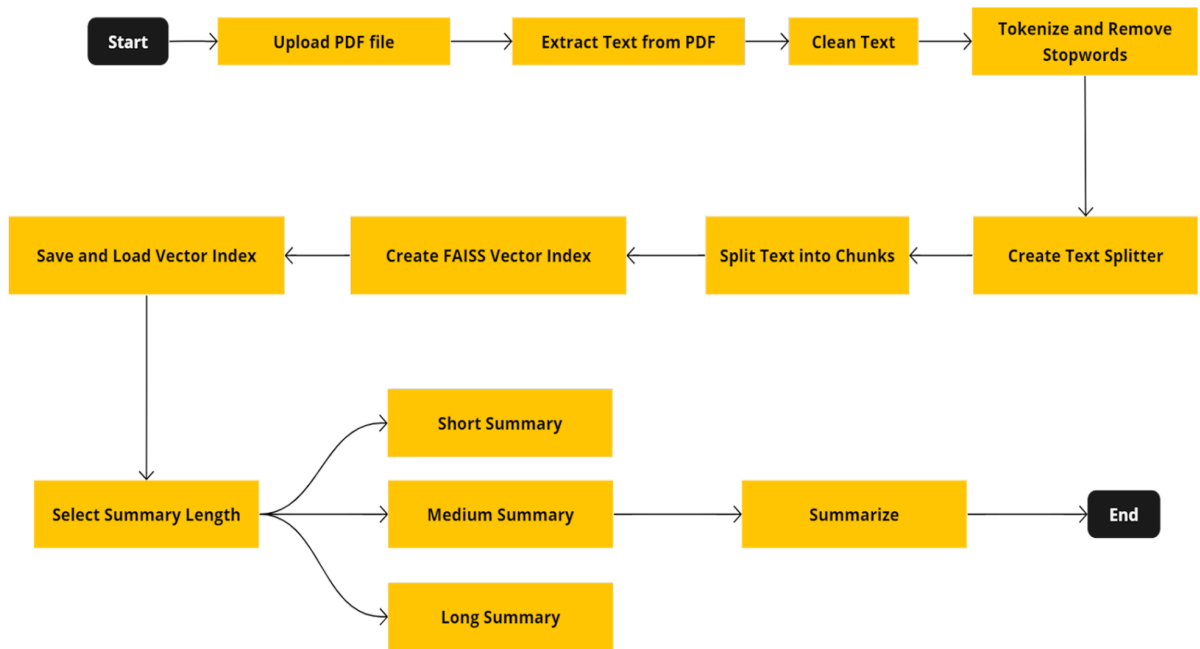
This vibrant word cloud represents the frequency of terms in a collection of research texts, with larger font sizes indicating higher occurrence. Central and prominent terms like "task," "data," "methods," and "performance" suggest a focus on empirical and methodological aspects of research. The word cloud offers a quick visual summary of key concepts and themes prevalent in the dataset, reflecting common language in academic studies or technical analyses.





**Figure 5: Workflow for automated text summarization**

The diagram presents a streamlined workflow for an automated text summarization system. It initiates with user interaction, where data is uploaded and sets into motion the subsequent processing stages. Text extraction follows, signifying the retrieval of textual content from the provided documents. This content is then subjected to NLP Processing, leveraging natural language techniques to refine and prepare the data for AI-based analysis. The heart of the system lies in the Summarization Using AI Models phase, where sophisticated algorithms distill the essence of the text into concise summaries. The process culminates with Summary Generation and Presentation, where these synthesized insights are neatly packaged and delivered back to the user. This cyclic workflow encapsulates the transformative journey of raw data into insightful, actionable summaries.



**Figure 6: Flowchart showcasing text summarization process**

This flowchart delineates the steps involved in a text summarization process, beginning with the user uploading a PDF file. The text is then extracted from the PDF and goes through a cleaning process to remove any unwanted characters or spaces. Subsequently, the text is tokenized and stopwords are removed to distill the content down to its most meaningful components. Following this, the text is split into chunks using a text splitter tool. These chunks are then indexed using a FAISS vector index, which is saved and loaded for retrieval. The user has the option to select the length of the summary they desire: short, medium, or long. Finally, the system generates the summary as per the selected length, concluding the process. This diagram succinctly captures the sequential steps and decision points involved in transforming a PDF document into a concise summary.

In conclusion, our extensive experiments underscored the effectiveness of our summarization tool. The results from the comparative analysis, ablation studies, and tuning experiments collectively affirmed the superiority of our approach. The visualization techniques not only enhanced our understanding of the model's internal workings but also served as a powerful tool to communicate these insights.

# Deployment

## 7.1 Introduction to Deployment

In this chapter, we document our journey to deploy our research paper summarization tool on Heroku, a cloud platform service. This account details the challenges we faced, the steps we took to overcome them, and the resources we utilized. Our initial attempt to deploy the application on PythonAnywhere was unsuccessful due to issues with python dependencies management. Consequently, we shifted our focus to Heroku.

## 7.2 Initial Steps and Challenges

### 7.2.1 Choosing Heroku

Our decision to switch to Heroku was driven by its robust support for Python applications and its seamless integration with Git. We anticipated a smoother deployment process, given Heroku's widespread usage and comprehensive documentation.

### 7.2.2 Preparing for Heroku Deployment

**Creating a New Branch:** We initiated our deployment process by creating a new branch in our GitHub repository named "chumbar/deploy-on-python-anywhere". This step was crucial as it allowed us to make necessary modifications tailored for Heroku's environment without affecting our main branch.

**Learning Curve:** Dedicating time to understand Heroku's deployment nuances was essential. Approximately two hours were spent studying Heroku's build and deployment process, which is intricately linked with Git.

## 7.3 Essential Files for Heroku

**Procfile Creation:** We created a 'Procfile', a crucial file for Heroku deployment, to specify the number of worker processes required for our application.

**Runtime Specification:** The 'runtime.txt' file was created to declare the specific Python runtime version.

## 7.4 Setting the Buildpack

An initial hurdle we encountered was with the buildpack configuration. Heroku's automatic buildpack selection erroneously opted for NodeJS, instead of Python. We resolved this by manually setting the buildpack to Python.

## 7.5 Overcoming Deployment Errors

### 7.5.1 Encountering H81 Errors

Our initial deployment attempts were met with H81 errors, indicating that we were pushing an empty application to Heroku. This required us to reevaluate and meticulously follow the necessary commands for a successful deployment.

### 7.5.2 Correct Command Execution

We executed a series of commands, ensuring proper login to Heroku, local testing of the app, and correct Git procedures:

*heroku login*

*source myenv/bin/Activate*

*pip install gunicorn*

*heroku create*

*heroku local web (for local testing)*

*git remote heroku main*

*git push heroku main*

However, despite these efforts, we still faced deployment challenges.

## 7.6 Branch-Specific Deployment

We discovered, from an obscure source, that to push from a different local branch to Heroku, a specific command format was necessary:

**git push heroku <local\_branch>:main.**

Implementing this command marked a turning point, successfully initiating the build process and displaying the frontend.

## **7.7 Addressing File Upload and Environment Variable Challenges**

### **7.7.1 Uploading Files Issue**

A significant challenge arose when we tried to upload files to the application. The issue stemmed from the fact that the placeholder values for critical variables, such as the PDF directory path and the `OPENAI_API_KEY`, were set to `'REPLACEME'`. This presented a dilemma because committing the actual values to Git would risk exposing our OPENAI API Key and disable it.

### **7.7.2 Implementing Heroku's Built-In Support for Secrets**

To navigate this challenge, we delved into Heroku's capabilities for managing secrets and environment variables. By setting the OPENAI API KEY and the PDF directory path as environment variables in Heroku, we could securely deploy our application without exposing sensitive data in our version control system.

## **7.8 Code Modifications for Directory Management**

### **7.8.1 Handling Non-Existent Directories**

Our next hurdle was related to the application's handling of non-existent directories. The original code was structured in a way that it would throw an error if the specified PDF directory did not exist. This was a significant impediment in the Heroku environment.

### **7.8.2 Code Change for Directory Creation**

We implemented a crucial code modification that altered the system's behavior. Instead of returning an exception when the PDF directory was not found, the updated code would now create the directory. This change was pivotal, as it allowed the website to function properly in the cloud environment, facilitating file uploads and processing.

## **7.9 Successful Deployment and Functionality Testing**

### **7.9.1 Final Deployment Steps**

After resolving the directory issue, we proceeded with the final deployment steps. We pushed our code to Heroku, ensuring all configurations and environment variables were correctly set.

### **7.9.2 Testing the Deployed Application**

We conducted thorough testing of the deployed application to validate its functionality. This included:

Testing file uploads and ensuring the PDF directory was correctly managed.

Verifying the integration with the OPENAI API for the summarization process.

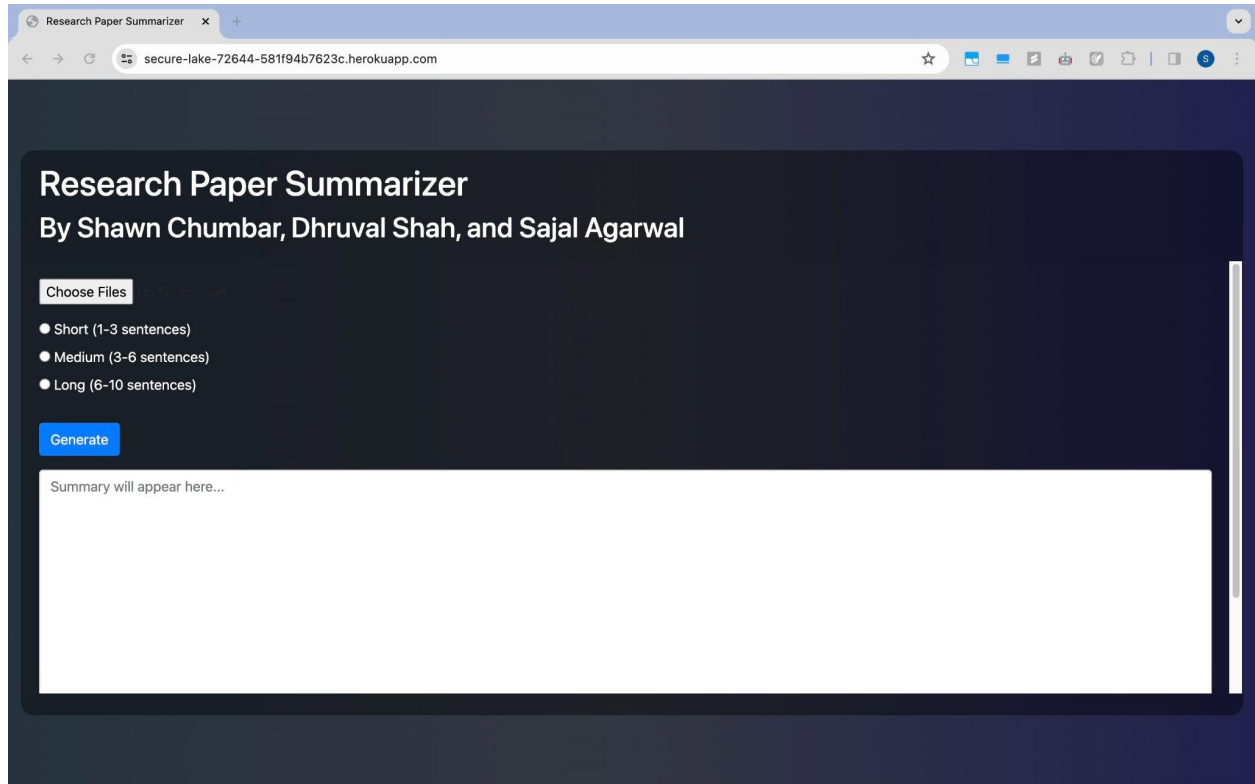
Assessing the application's responsiveness and performance on the Heroku platform.

### **7.10 Resolving Remaining Issues**

During testing, we addressed minor issues related to performance and user interface, fine-tuning the application to ensure a smooth user experience.

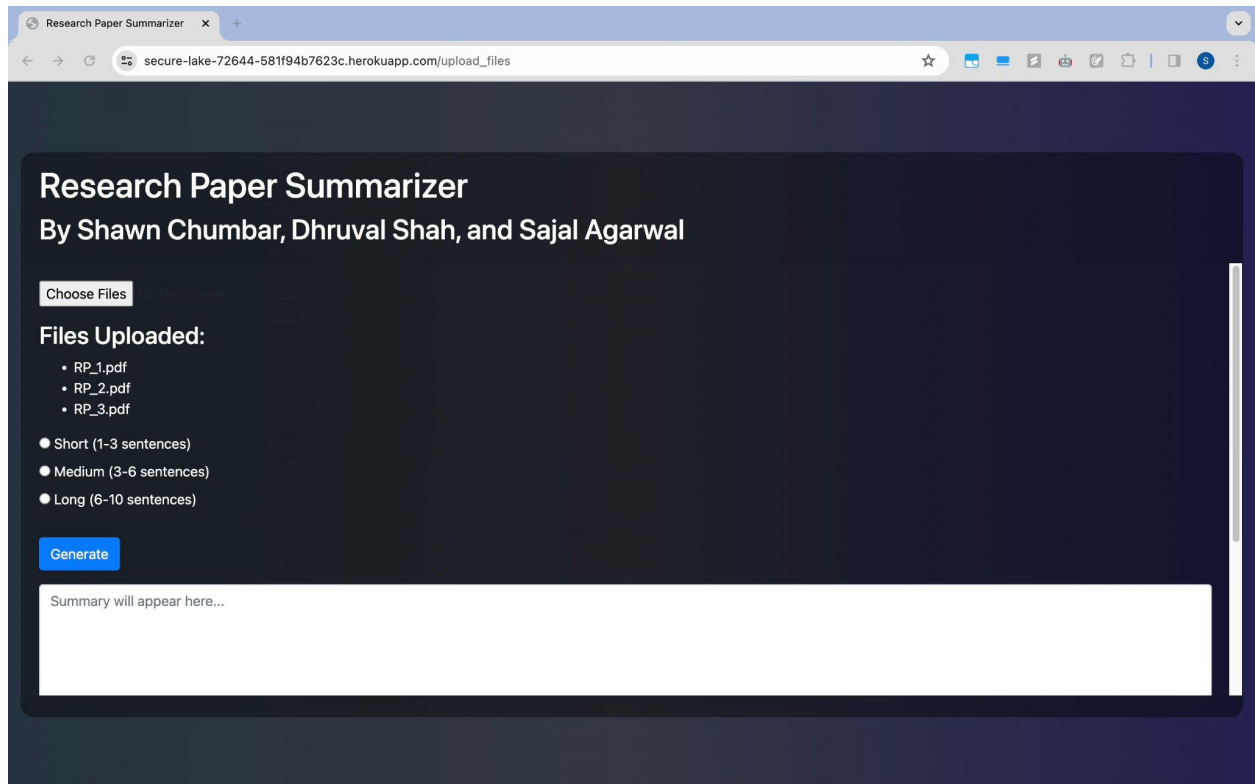
### **7.11 Application Screenshots**

After deployment of our application, we were able to confirm that the front end was functioning correctly. Please see below for figures related to the application frontend.



**Figure 7: Initial view of the frontend application.**

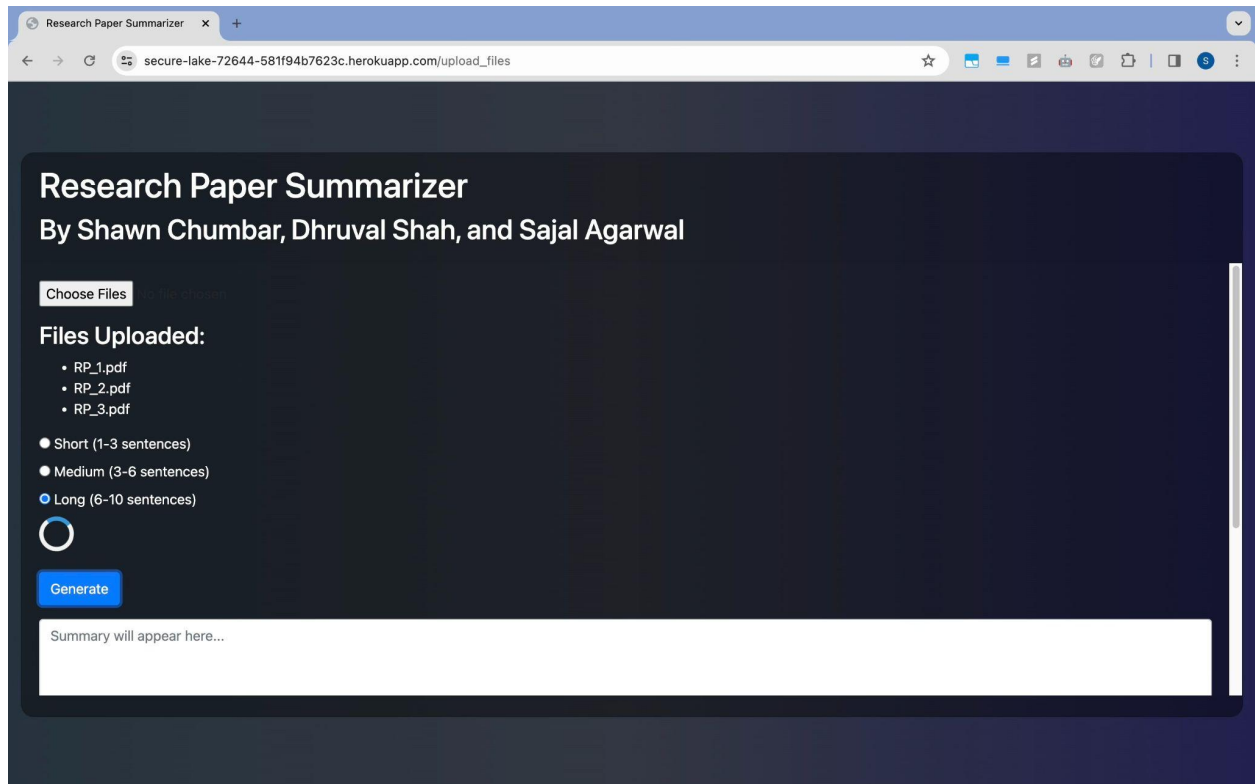
The figure above shows the initial view of the front end application. In the user interface, we have a “Choose Files” button to select what research papers to upload. We also have radio buttons that specify the length of the summary (i.e. Short, Medium, and Long). Each summary length has a certain number of sentences tied to it (i.e. Short is 1-3 sentence summary, Medium is 3-6 sentence summary, and Long is 6-10 sentence summary). The default length is 15 sentences when none of the radio buttons have been selected.



**Figure 8: View of frontend application after research paper PDF files.**

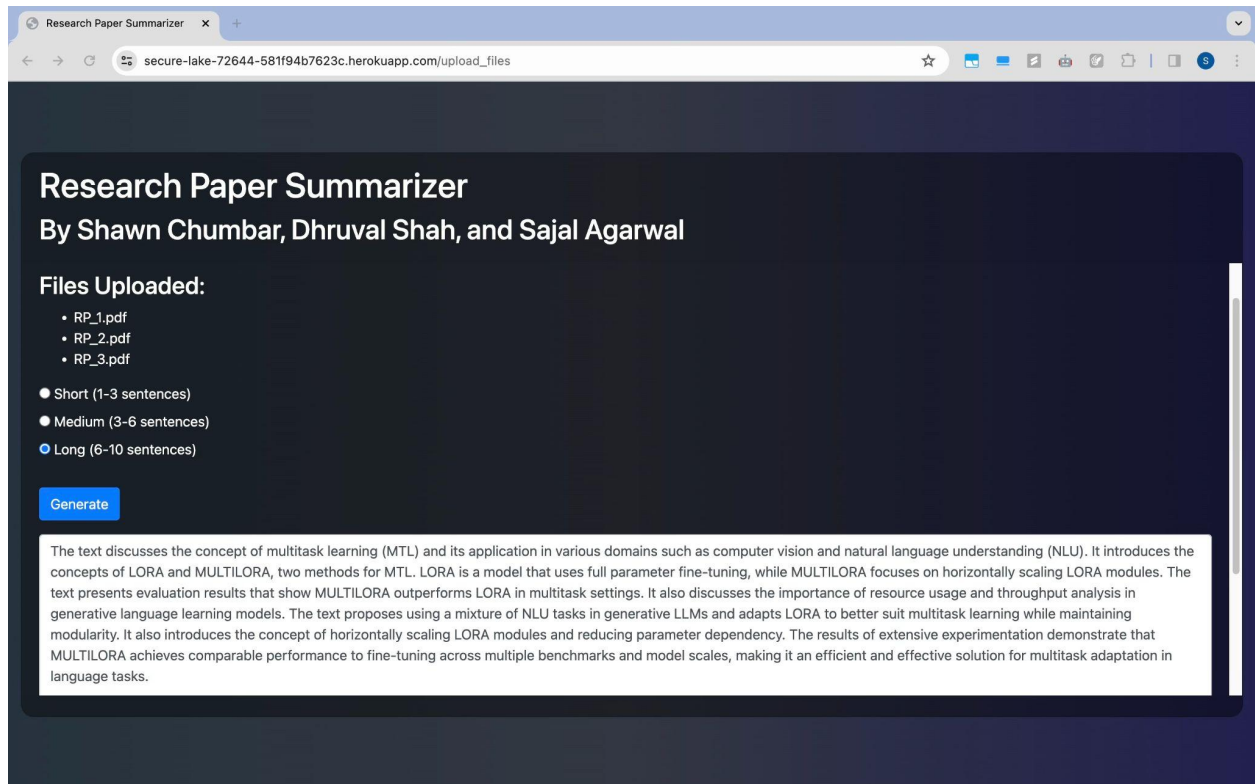
After uploading the research papers, you will see that there is a new section called “Files Uploaded”. Within this section, the list of file names that we have uploaded are shown. This ensures that the user is uploading the correct files.





**Figure 9: View of frontend application after clicking on generate button.**

After clicking on the “Generate” button, a loading spinner is shown to the user indicating that the model has begun working. In the backend, this includes doing various tasks such as deleting any files from previous summaries, extracting the text from the PDF files that were uploaded, cleaning and tokenizing the text, and using embeddings and vector index to create a summary of the text.



**Figure 10: Frontend application after loading research paper summary.**

After the summary for the attached research papers is generated, the summary will be displayed in the text section of the frontend. In the above figure, you can see that the summary was generated.

## 7.12 Conclusion

The deployment of our research paper summarization tool on Heroku was a journey replete with challenges and learning opportunities. Through perseverance and resourceful problem-solving, we successfully deployed a functional and secure application in the cloud. This chapter not only documents our deployment process but also serves as a testament to our team's ability to navigate complex cloud environments and adapt to unforeseen challenges.

## Conclusion

The project on summarizing research papers has been a fascinating journey, revealing much about the potential and challenges of working with complex academic texts. It's about making the knowledge more accessible and usable, not just for a select few but for everyone who seeks to learn and grow. The potential applications and future extensions of this technology are as boundless as the realms of knowledge it seeks to summarize. Here's a closer look at the key takeaways and a glimpse into what the future might hold for this exciting area:

1. **Robust Summarization Capabilities:** The system's ability to condense lengthy, complex papers into brief, understandable summaries is one of the most important results. This feature is a significant advancement in the accessibility of academic knowledge, particularly for those who might not have the time or background to read lengthy articles.
2. **Versatility Across Topics:** The model shows great potential as an adaptable tool in academic research because of its capacity to handle a wide range of topics, from the complex nature of neural networks to the specifics of multi-task learning. In an environment with plenty of resource range, this broad application is critical.
3. **Advancements in Information Retrieval:** Utilizing FAISS for vector index creation has enhanced the process of searching and retrieving relevant information from large text datasets. This method has shown efficiency and accuracy, crucial in dealing with extensive academic literature.
4. **Quality of AI-Generated Contextual Understanding:** Another key learning was the AI's ability to not just summarize text, but to do so with a nuanced understanding of context and content. The model demonstrated an impressive capability to grasp and convey complex academic concepts, theories, and methodologies. This level of contextual understanding and interpretation is a significant step forward in AI-driven text analysis, particularly in the field of academic research where accuracy and depth of understanding are paramount. It highlights the advanced state of current NLP technologies and their potential to further evolve in handling specialized, context-heavy content.
5. **Navigating Complex Academic Language:** The project also highlighted the complexities involved in processing academic language, which often includes specialized terms and elaborate structures. Overcoming these challenges was key to achieving accurate summarization, underscoring the sophistication of the employed NLP techniques.

## Future Extensions and Applications

For future directions, the project can be expanded in several promising ways:

1. **Scalability and Performance Optimization:** A key objective will be to enhance the system's capacity to manage a greater volume of data requests. Optimizing the underlying algorithms and infrastructure to deliver rapid summaries is critical, especially when scaling up to process an extensive array of research papers across various disciplines. This will involve refining the model's efficiency and ensuring robustness under heavy loads, making it a more practical tool for institutions with large databases of academic work.
2. **Multi-Format Documents:** The project aims to extend beyond PDFs to accommodate multiple document formats, such as HTML, DOCX, and LaTeX, which are common in academic research. This will enable the system to process and summarize a broader spectrum of research materials, making the technology more versatile for different data science applications where multiple formats are prevalent.
3. **Interactive Summarization:** Incorporating interactivity into the summarization process is another exciting avenue. Users could engage with the system by asking questions or requesting more information on specific segments of the research paper. The AI would respond with targeted summaries or detailed expansions, tailoring the information to the user's immediate needs. This could transform the summarization tool into an interactive learning platform, enhancing the user experience and providing a richer, more dynamic engagement with the research material.
4. **Broader Academic Scope:** Expanding the system to encompass a wider range of disciplines, including the humanities and social sciences, could significantly enhance its utility. This would test and potentially improve the model's ability to adapt to various styles of academic writing and content.
5. **Multilingual Expansion:** Extending the model to translate and summarize research papers in multiple languages would be a significant stride in making global academic knowledge universally accessible. This could bridge language barriers in academia, fostering cross-cultural research collaboration.
6. **Customizable Summaries:** Developing the capability for users to specify their areas of interest for more personalized summaries could be a game-changer. This feature would tailor the output to individual needs, increasing the utility of the summaries.
7. **Integration with Digital Libraries:** Implementing this summarization technology in academic search engines or digital libraries could revolutionize how researchers and students interact with academic content. It would enable quick and efficient

browsing of large volumes of research, helping users identify relevant papers more effectively.

In essence, this project has not only achieved its immediate objective of effectively summarizing research papers but has also opened exciting possibilities for further advancements in academic text processing. The potential applications and extensions of this technology are vast and hold the promise of transforming how academic content is accessed, understood, and utilized.